

# Salamander: computer vision for wildlife research

Christopher Patton

March 5, 2013

## Abstract

This document outlines an image processing application called Salamander. Salamander is comprised of a set of algorithms for automatically detecting and tracking targets of interest in a video stream. It is particularly useful for filtering and segmenting video produced by field cameras for wildlife and behavioral research.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Algorithms</b>	<b>2</b>
2.1	Image processing . . . . .	2
2.2	Detection . . . . .	3
2.3	Tracking . . . . .	3
2.4	Segmentation . . . . .	3
<b>3</b>	<b>The Salamander toolkit</b>	<b>3</b>
3.1	Build and install . . . . .	4
3.2	A quick example . . . . .	4
3.3	Usage . . . . .	4
3.3.1	<code>binthresh</code> and <code>binmorph</code> . . . . .	5
3.3.2	<code>filter</code> . . . . .	5
3.3.3	<code>detect</code> . . . . .	5
3.3.4	<code>segment</code> . . . . .	6
3.4	Processing a feed in realtime . . . . .	6
3.5	Code documentation . . . . .	6
<b>4</b>	<b>Machine learning</b>	<b>6</b>

# 1 Introduction

[TODO] Overview of approach. Define each task: detection, tracking, and segmentation.

## 2 Algorithms

I now describe our approach to target detection and tracking in and segmentation of a video stream. In essence, this is an application of computer vision in which we make some basic assumptions about the video and our targets:

1. The camera is fixed.
2. The rate at which pictures are taken is fixed.<sup>1</sup>
3. The targets are roughly the same size and color; i.e., they are the same species (not looking for deer and mice at the same time) and roughly the same distance away from the camera when we detect them.

The reasons for these assumptions will become apparent in this section. They are necessary for the basic functionality of the algorithms we'll describe, but it may be possible to disregard at least (3) with machine learning. We'll explore this possibility in section 4. However, we must assume (1) in any case.

Initial image processing, detection, tracking, and segmentation are layered. By this we mean that in order to achieve one task, all preceding tasks must be accomplished first. We describe each layer and the corresponding algorithms and data structures in turn. In the following sections, a video stream is represented by a sequence of video frames:

$$\langle v_0, v_1, v_2, \dots, v_{n-1}, v_n \rangle$$

### 2.1 Image processing

The first step in our approach is to calculate the pixelwise difference of two adjacent frames in a video stream. The resulting image, which we'll refer to as the *delta* image, tells us which regions of the image changed within that unit of time. We'll perform some additional filtering, then search the image for *blobs*, contiguous regions of *delta* that correspond to significant changes. At this point, image processing is finished and we can proceed to detection.

**delta**  $\delta(v_{i-1}, v_i) \rightarrow d_i$ . Input images are typically JPEG format. Each pixel of a JPEG image is represented by three bytes corresponding to the red, blue, and green luminance. To simplify the following steps of the pipeline, we first convert  $v_{i-1}$  and  $v_i$  to gray scale images, whose pixels are floating point

---

<sup>1</sup>It may be necessary in certain implementations to overcome this assumption. For example, a field camera may transmit its video stream over a network. The rate at which pictures are taken can be set to depend on network conditions. However, It is possible to assume a lowest-possible rate and tune Salamander accordingly.

numbers (since we converted from bytes). Next we subtract  $v_i$  from  $v_{i-1}$  and store the absolute value of each pixel. Now, all changes in the frame within the time points  $i - 1$  to  $i$  are reflected in the delta image  $d_i$ . We can analyse these changes in order to determine if a target has appeared in frame.

**binary threshold**  $b(d_i) \rightarrow b_i$ . The next step converts all pixels to one of two values {white, black}, which is needed for subsequent filters. This step also allows us to deal with noise artifacts related to the quality of the image taken, e.g. blurriness, compressing and decompressing JPEGs, gray scale filtering, etc. More to the point, if a pixel has changed just a little, we set it to black; if it's changed a lot, we set it to white. For each pixel  $p$  in  $d_i$  and threshold range  $l$  and  $h$  where  $l < h$ ,

$$b_i(p) = \begin{cases} \text{white} & l \leq d_i(p) \leq h \\ \text{black} & \text{otherwise.} \end{cases}$$

[TDOD] experiment with upper threshold higher than 255? Since floating points in ImageType can be higher I'm pretty sure.

**binary morphology**  $m(b_i) \rightarrow m_i$ . finally, the last filter enables us to deal with real noise. [TODO]

## 2.2 Detection

connectedComponents(im). How to handle noise.

## 2.3 Tracking

Shifting over merged blobs. Multiple targets.

## 2.4 Segmentation

Data structures. Testing if gap has target.

# 3 The Salamander toolkit

Salamander is a C++ extension to the Insight Registration and Segmentation Toolkit<sup>2</sup>. ITK provides a comprehensive array of image processing algorithms and is designed to be useful in many applications. A key design feature of ITK is its support of generic programming. As well as processing two-dimensional images, it can process multidimensional image and mesh data, making it a successful framework for medical image analysis. This architecture is extremely useful when, for example, a project must adapt and scale to new data formats. However, in our computer vision application, we've restricted ourselves to two dimensional images and affectively ignore the vast majority of ITK's features. As a result, we encurr the overhead of generic programming without gaining its benifits.

---

<sup>2</sup>ITK is an open source product by Kitware and is available at <http://www.itk.org>.

Therefore, it may be prudent to switch to a framework that is centered around computer vision. We will consider the feasibility of using OpenCV<sup>3</sup> instead of ITK. OpenCV is written in C and doesn't support generic programming. As long as the image processing algorithms we need (pixelwise subtraction, binary thresholding and morphology, connected component analysis) are comparably efficient, OpenCV will be a viable replacement.

Another advantage of OpenCV over ITK is that it provides a well-known means of implementing machine learning, something that we'll explore in future iterations of the Salamander project.

### 3.1 Build and install

Both ITK and Salamander should be built from source. ITK is available for download at <http://www.itk.org/ITK/resources/software.html>. As of this writing, Salamander is hosted on github at <http://www.github.com/cjpatton/salamander>. The build process for both is identical, as they use `cmake`. In the source directory, under GNU/Linux:

```
$ mkdir build && cd build
$ cmake ../
$ make
$ sudo make install
$ sudo ldconfig
```

ITK has a large amount of code and will require a couple hours to finish. As always, two cores can be assigned to building with `make -j 3`.

[TODO] Windows instructions.

### 3.2 A quick example

If you've successfully installed Salamander, you can try this tutorial to see how it works. In the source directory, go to `ex/three`. This footage is from an experiment with mice, in which they are the choice of two environments. The goal was to find out which they prefer. Type the following at the terminal:

```
$ ls video*.jpg | segment -t 20 60 -m 1 15
```

The program will output the chunks of video in which a mouse appears in frame. Also, a series of images are produced called `tracking*.jpg`, in which `segment` has drawn bounding boxes around the location of the target.

### 3.3 Usage

After installation, there are a number of top-level programs available. This section documents what they do and how to use each of them. Each Salamander program accepts a list of filenames on standard input and sorts them

---

<sup>3</sup><http://opencv.org>.

alphanumerically to ensure they are processed in the correct order. Except for **binthresh** and **binmorph**, all of the programs have the same command line options which enable tuning of the image processing parameters. These are:

- t *l, h* Binary threshold range, where  $0 \leq l \leq h \leq 255$ . Any pixel whose value falls between *l* and *h* become white: otherwise, black. Although we can make adjustments to reduce noise here, this value should be set so that an interesting blob appears white after this stage. Therefore, a wide filter is commonly used, such as *l* = 20 and *h* = 60.
- m *e, d* Binary morphology factors. *e* specifies the number of erosion iterations and *d* specifies the number of dilation rounds. Typically,  $e < d$ . There are two considerations for this parameter: one, the relative size of the target, and two, the frequency of noise.
- s *s* Image shrink factor, where  $s > 0$ . Image processing is far and away the bottleneck of Salamander. We can increase throughput by shrinking the image. The trade off is that we can lose information and, potentially, miss small targets. If  $s = 2$ , then the image is shrunk by a factor of two. Default is  $s = 1$  (no shrinking).
- f **prefix** Output file prefix. Some programs output files.

### 3.3.1 **binthresh** and **binmorph**

These simple programs apply image processing filters to the images and output the result as a sequence of images. They are useful for experimenting with parameters on a small set of images. **binthresh** accepts two arguments corresponding to the binary threshold range. If none are provided, only the image difference filter is applied. **binmorph** accepts four arguments corresponding to the morphology factors and the threshold range. If two are provided, these are assumed to be the morphology factors and the threshold defaults to  $l = 20$  and  $h = 60$ .

E.g.: `$ ls video*.jpg | binmorph 1 10 20 60`

### 3.3.2 **filter**

This program is functionally similar to **binmorph** except that it provides all the normal tuning parameters and provides more useful information. In particular, for each delta image, it outputs a list of blobs. **filter** is useful for testing parameters over a large set of images. In the future, it will be modified to be useful for collecting statistics [IMPLEMENT].

E.g.: `$ ls video*.jpg | filter -t 20 60 -m 20 60`

### 3.3.3 **detect**

Currently, this program is being used to test the more advanced detection scheme discussed in 2.2. [IMPLEMENT] Once this is finished, this program will simply output the images that contain targets.

### 3.3.4 `segment`

Currently, `segment` is the highest level program. It applies the image processing filters, detects the presence of targets, tracks the target in time, and outputs segments of video that contain targets. [TODO] Rename files by segment number, dump to disk. Delete all files that don't appear in a chunk.

E.g.: `$ ls video*.jpg | segment -s 3 -t 20 60 -m 20 60`

## 3.4 Processing a feed in realtime

[NOT IMPLEMENTED] This feature should be implemented as a flag at the command line.

## 3.5 Code documentation

[TODO] I hope to generate documentation with doxygen. This will be plain HTML.

# 4 Machine learning

[FUTURE] I still need to research how and if to do this.